BlueCloud

DATABRICKS AI/ML ENGINEER

ONSHORE/NEARSHORE — REMOTE

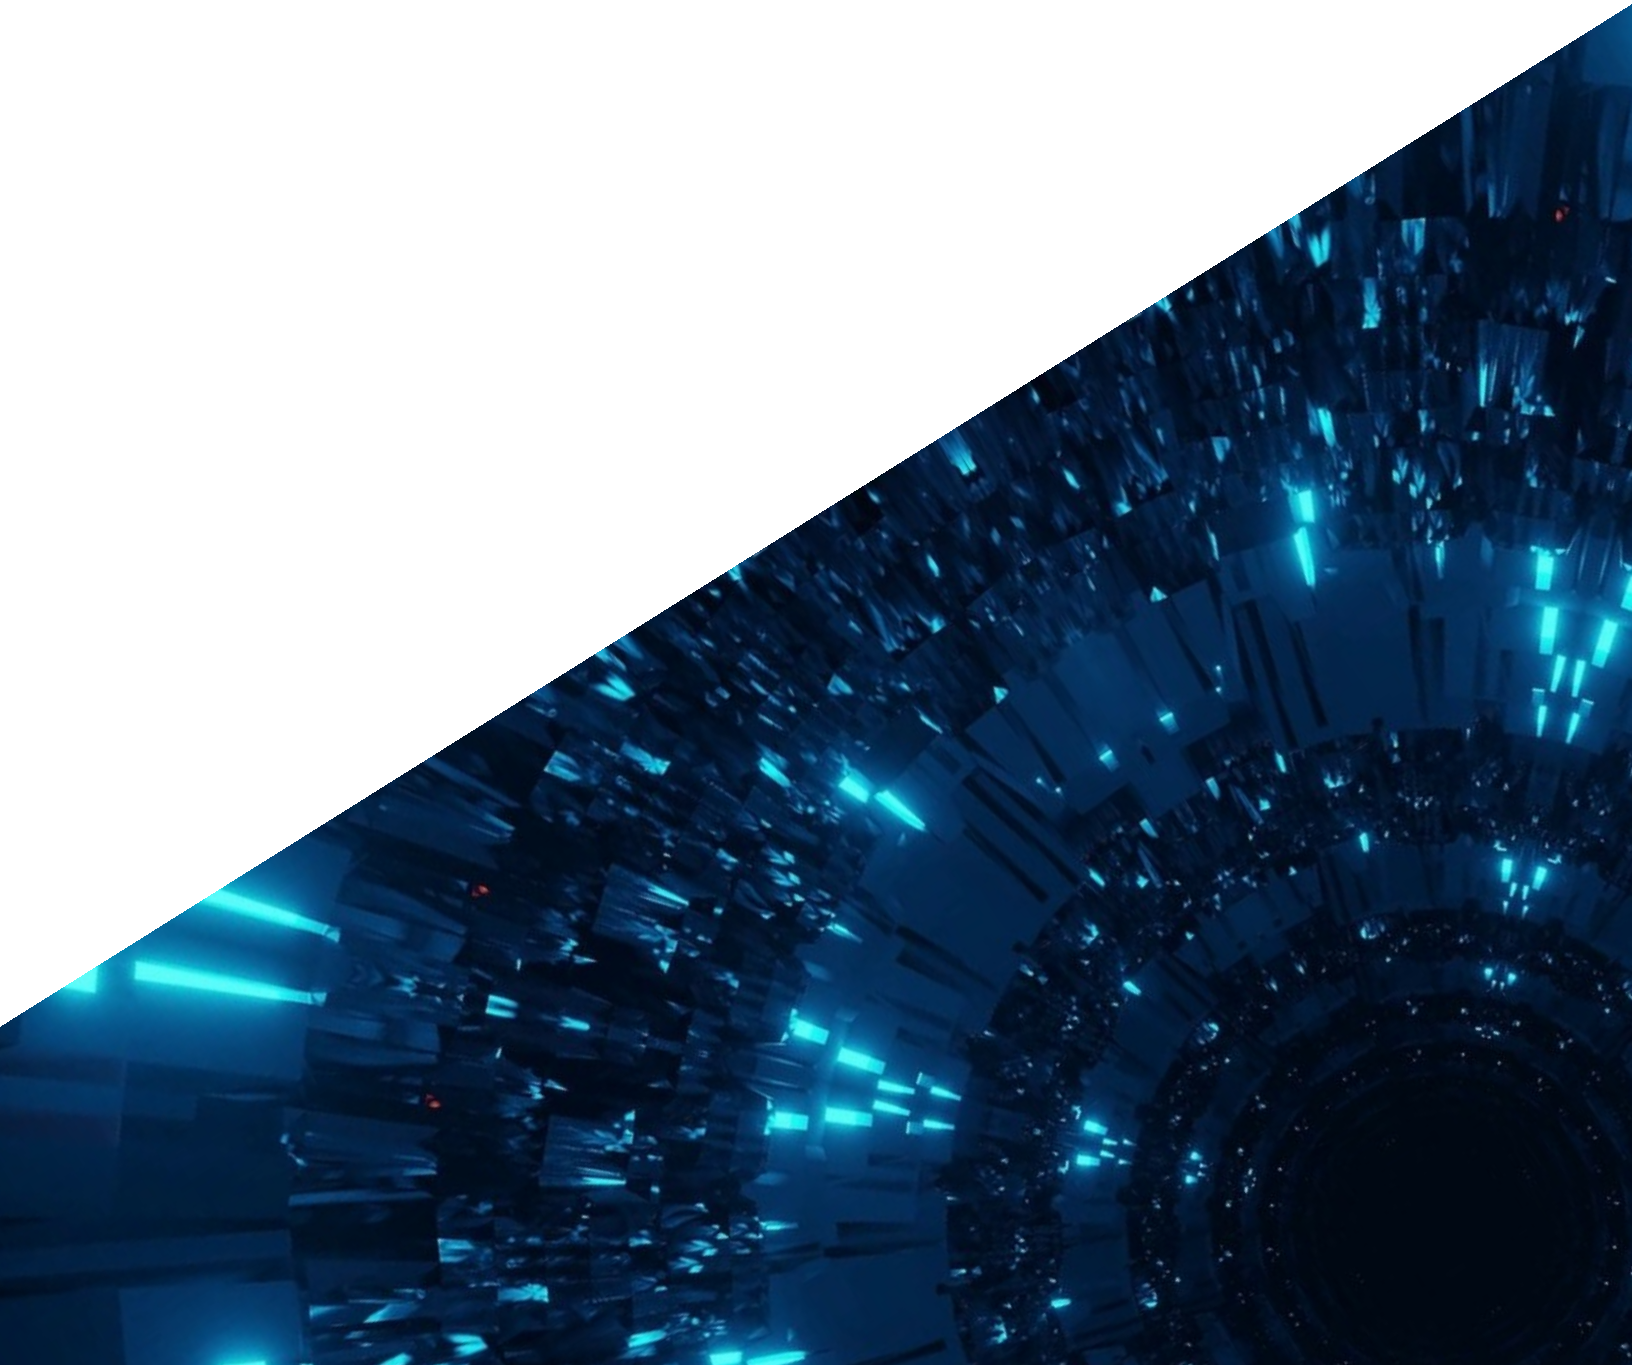

ABOUT US

Our firm's name says a lot. Simply put, Blue.cloud enables enterprises to transform to cloud. As a Cloud-Only / Cloud-Forward transformation company, we help organizations to enable and speed their journey to cloud.

Blue.cloud recognizes the cloud-only nature of our digital transformation company that specializes solely in implementing leading cloud technologies and solutions for our clients with data & analytics experiences from 2004.

Blue.cloud's innovative approaches and tailored solutions designed to meet our customers' unique requirements are the core components of our service offerings. With a balanced view of strategy, people, process and technology, we work with companies to understand their specific business needs to develop and implement practical business strategies and technology solutions.

By leveraging our strategic partnerships with cloud superpowers such as Microsoft Azure, Google Cloud, AWS, Snowflake & ThoughtSpot, Blue.cloud continues to highlight the opportunity set at hand for cloud migration, often consulting on cloud strategies and subsequently performing the implementation. With our extensive experience in digital-focused services, Blue.cloud is well-positioned to help our customers take advantage of the upcoming cloud-centric investment cycle.

Join Blue.Cloud: a global leader in cloud solutions with an exceptional focus on innovation, customer service, and employee engagement.

DATABRICKS AI/ML ENGINEER

POSITION

Databricks AI/ML Engineer

LOCATION

Onshore/Nearshore — Remote

Responsibilities

- Design, develop, and deploy machine learning models on the Databricks platform.
- Build and operationalize end-to-end ML pipelines using Databricks, including feature engineering, model training, evaluation, and production deployment.
- Leverage Databricks MLflow for experiment tracking, model registry, and end-to-end ML lifecycle management.
- Build agentic AI solutions on Databricks to support decision-making use cases, including rule-based alerts, anomaly detection, and automated insights delivery for business managers.
- Collaborate with data engineering teams to design and implement scalable, high-performance data architectures on Databricks Lakehouse Platform.
- Work with large-scale structured and unstructured datasets using Databricks Delta Lake and Apache Spark.

Qualifications

- Experience with Databricks AI/BI Genie for conversational analytics, self-service insights, and business intelligence use cases.
- Hands-on experience with Databricks Agent Bricks for building, orchestrating, and operationalizing AI agents.
- Strong understanding and implementation experience with Generative AI solutions on Databricks, including RAG (Retrieval-Augmented Generation) and LLM-based applications.
- A good understanding of retail and CPG data domains is preferable.